

◇ DATA ◇ CODE ◇ MODELS ◇ INFRA

MLOps Best Practices

Shipping machine learning systems that stay reliable, reproducible, and maintainable — long after the demo.

Adapted from an article by Yadidiah Kanaparthi

Use ← → or Space to navigate_

QUIZ · MLFLOW

QUESTION 1 OF 10

In MLflow, which of these would you log as a Parameter rather than a Metric?

- A The learning rate used for training ✓
- B The model's test accuracy
- C A confusion matrix image
- D The trained model file

Parameters are inputs you fix before training starts — a learning rate is exactly that.

QUIZ · MLFLOW

QUESTION 2 OF 10

What is the correct relationship between an MLflow Experiment and a Run?

- A A Run can belong to multiple Experiments at once
- B An Experiment groups together multiple Runs of the same project ✓
- C A Run and an Experiment are interchangeable terms
- D Experiments are created only after all Runs finish

An Experiment is the folder; each training run you log inside it is one Run.

QUIZ · MLFLOW

QUESTION 3 OF 10

Which MLflow function is used to save a trained model as a reusable, versioned entity in the Model Registry?

A `mlflow.log_metric()`

B `mlflow.log_artifact()`

C `mlflow.register_model()` (or `log_model(..., registered_model_name=...)`)



D `mlflow.set_experiment()`

`register_model()` (or `log_model` with `registered_model_name`) is what actually creates a versioned entry in the Model Registry.

QUIZ · MLFLOW

QUESTION 4 OF 10

Why is it useful to log the same metric (e.g. F1-score) across several Runs with different hyperparameters?

- A It isn't useful — one Run is always enough
- B It reduces the model's training time
- C MLflow requires at least 3 Runs per Experiment to function
- D It lets you compare configurations objectively instead of relying on memory ✓

Logging the same metric across configurations is what turns tuning into a comparison instead of a guess.

QUIZ · MLFLOW

QUESTION 5 OF 10

What's the main difference between what Scikit-learn does and what an experiment-tracking tool like MLflow does?

- A Scikit-learn performs the actual training/prediction; MLflow records what happened during that process ✓
- B Scikit-learn logs metadata; MLflow trains the model
- C They do the same job, just with different syntax
- D MLflow replaces the need for a machine learning library entirely

Scikit-learn does the actual modeling work; MLflow just records what happened so you can find it again later.

QUIZ · PREPROCESSING & ENSEMBLE METHODS

QUESTION 6 OF 10

When treating outliers with the IQR (interquartile range) method, "capping" (winsorizing) a value means:

A Deleting the row that contains it

C Replacing it with the column mean

B Replacing it with the nearest boundary of the acceptable range instead of removing the row ✓

D Ignoring it during training but keeping it for evaluation

Capping clips the value to the nearest acceptable boundary instead of throwing the row away.

QUIZ · PREPROCESSING & ENSEMBLE METHODS

QUESTION 7 OF 10

Why might engineering a ratio between two correlated numeric features (instead of using them separately) help a classifier?

- A It always improves accuracy regardless of the model
- B It reduces the dataset size
- C It can capture a relationship between the two variables that's more directly relevant to the target than either raw value alone ✓
- D It removes the need for a train/test split

A ratio can encode the relationship between two variables in a way neither raw feature captures alone.

QUIZ · PREPROCESSING & ENSEMBLE METHODS

QUESTION 8 OF 10

Why is it risky to reuse a single already-fitted ColumnTransformer/preprocessor object across two different model pipelines?

- A It isn't risky — this is the recommended approach
- B It causes a memory leak that crashes the kernel
- C ColumnTransformer objects can only be used once, ever, and must be deleted after
- D Fitting it again inside the second pipeline can silently change its state, affecting the first pipeline too — use clone() to avoid this ✓

Fitting a shared transformer inside a second pipeline can quietly mutate it — clone() keeps the two independent.

QUIZ · PREPROCESSING & ENSEMBLE METHODS

QUESTION 9 OF 10

What does `handle_unknown='ignore'` do when passed to a `OneHotEncoder`?

- A Prevents an error when the encoder sees a category at prediction time that it never saw during training ✓
- B Ignores missing values in the training set
- C Skips encoding for columns with too many unique values
- D Disables one-hot encoding and falls back to label encoding

It tells the encoder to zero-fill an unseen category at inference instead of raising an error.

QUIZ · PREPROCESSING & ENSEMBLE METHODS

QUESTION 10 OF 10

Conceptually, how does Boosting (e.g. XGBoost) differ from Bagging (e.g. Random Forest)?

- A Boosting trains trees independently in parallel; Bagging trains them sequentially
- B Boosting trains trees sequentially, each one correcting the errors of the previous ones; Bagging trains trees independently on bootstrap samples ✓
- C Boosting only supports regression tasks
- D There's no real difference between the two approaches

Boosting corrects errors sequentially, tree after tree; Bagging averages independent trees trained in parallel.

01 · FOUNDATIONS

What Is MLOps?

MLOps applies the discipline of DevOps — automation, testing, versioning, monitoring — to the unique challenges of machine learning: data that changes, models that decay, and experiments that must be reproducible.

FIG. 01

DevOps

Continuous integration, continuous delivery, infrastructure as code, automated testing.

FIG. 02

Data Engineering

Pipelines, quality checks, versioning, and governance for the data feeding every model.

FIG. 03

Machine Learning

Experimentation, training, evaluation, and the statistical nature of model behavior.

MLOps vs. DevOps: What's Different

MLOps grew out of DevOps — but machine learning introduces failure modes software engineering doesn't have.

Shared Principles

- Process automation
- Continuous integration & deployment
- Collaboration & communication
- Scalability & reliability
- Monitoring & feedback loops

Unique to MLOps

Reproducibility

Same code + same data can still give different results — version data, seeds, hyperparameters, and environment.

Scaling Training & Inference

Training needs specialized hardware (GPUs); serving must stay fast and cheap enough to be worth it.

Model Drift

Accuracy degrades as real-world data shifts away from what the model was trained on.

Team Composition

Researchers and production engineers have different skills and paces — they have to learn to work together.

01 · FOUNDATIONS

Why It Matters

A notebook that produces a good metric is not a product. Most ML initiatives stall between prototype and production — not because the model is wrong, but because nothing around it is operationalized.



FIG. 01

The prototype-to-production gap

Notebooks don't version data, don't handle failure, and don't tell you when they're wrong in production.

FIG. 03

Everything drifts

The world the model was trained on keeps changing after deployment — user behavior, upstream data, seasonality.

FIG. 02

Silent failure

Unlike a crashing service, a model can keep returning confident, plausible, and wrong answers indefinitely.

FIG. 04


Reproducibility debt

Without versioned data, code, and config, "it worked on my machine" becomes "it worked once, three months ago."

01 · FOUNDATIONS

The ML Lifecycle Is a Loop

Unlike traditional software, the lifecycle never really ends at "deploy." Production feedback continuously reshapes the next iteration.




↻ Retraining feeds back into Collect & Prepare

02 · REPRODUCIBILITY

Version Everything

If you can't reconstruct exactly what produced a model, you can't debug it, audit it, or trust it.
Four things need a version history.



02 · REPRODUCIBILITY

Data Versioning & Lineage

Code review answers "what changed and why." Data needs the same answer — which snapshot, transformed how, by which job.

FIG. 01

Version like code

Tools such as DVC, LakeFS, or Delta Lake let you tag, diff, and roll back datasets the way Git handles source files.

FIG. 02

Track lineage

Record which raw sources, transformations, and jobs produced each downstream table or feature.

FIG. 03

Debug backwards

When a model misbehaves, lineage lets you walk back to the exact upstream change that caused it.

FIG. 04

Enable audits

Regulated industries need to prove exactly what data trained a production model, months later.

02 · REPRODUCIBILITY

DVC in Practice

Version datasets and pipelines alongside Git — reproducible by construction.

TERMINAL

```
# version a dataset
dvc add data/train.csv
git add data/train.csv.dvc
git commit -m "Add training data v1"

# run the full pipeline
dvc repro
```

DVC.YAML – PIPELINE DEFINITION

```
stages:
  preprocess:
    cmd: python src/data/preprocessing.py
    deps:
      - src/data/preprocessing.py
      - data/raw/transactions.csv
    outs:
      - data/processed/clean.csv
  train:
    cmd: python src/models/train.py
    deps:
      - src/models/train.py
      - data/processed/clean.csv
    metrics:
      - metrics.json
```

Every dataset version is tied to a Git commit – checkout any commit and dvc checkout restores the exact data that produced it.

02 · REPRODUCIBILITY

Feature Stores

The most common production bug in ML: features computed differently in training than in serving — "training/serving skew."

FIG. 01

Offline store

Historical, point-in-time correct features used to build training datasets.

FIG. 02

Online store

Low-latency key-value lookups serving the same feature definitions at inference time.

Why it matters

A feature store defines each feature's logic once and reuses it in both paths — eliminating an entire class of silent production bugs.

03 · EXPERIMENTATION

Track Every Experiment

If a result isn't logged, it didn't happen. Experiment tracking turns ad-hoc notebook runs into a searchable, comparable history.

FIG. 01

Parameters

Hyperparameters, feature sets, random seeds — the recipe for each run.

FIG. 02

Metrics

Loss curves, accuracy, business metrics — logged per step, not just the final number.

FIG. 03

Artifacts

Model weights, plots, and evaluation reports attached to the run that produced them.

Tools in this space (e.g. MLflow, Weights & Biases) exist mainly to make "which run was that?" a solved problem.

03 · EXPERIMENTATION

MLflow in Action: Track Everything

A config-driven training function logs params, metrics, and the model itself on every run.

```
def train_model(config_path):  
    config = yaml.safe_load(open(config_path))  
    mlflow.set_experiment(config['experiment_name'])  
  
    with mlflow.start_run(run_name=config['run_name']):  
        mlflow.log_params(config['model_params'])  
  
        X_train, y_train = load_data(config['data_path'])  
        model = RandomForestClassifier(**config['model_params'])  
        model.fit(X_train, y_train)  
  
        mlflow.log_metric("accuracy",  
                          accuracy_score(y_val, model.predict(X_val)))  
        mlflow.sklearn.log_model(model, "model")
```

Compare experiments visually in the MLflow UI

Reproduce any run from its logged params + code version

Deploy models directly from the registry

03 · EXPERIMENTATION

Reproducible Environments

"Works on my machine" is not acceptable for a training job that needs to run identically today, next month, and on a teammate's laptop.

FIG. 01

Containers

Docker images pin the OS, drivers, and system libraries around the training/serving code.

FIG. 02

Locked dependencies

Exact, hashed package versions — not loose ranges — so a rebuild resolves identically.

FIG. 03

Infra as code

Clusters, GPUs, and networking defined in version-controlled templates, not clicked together by hand.

03 · EXPERIMENTATION

How to format your Ready-Production Project

Stop using flat notebooks, A production-ready project separates concerns — data, source code, configs, and tests each get their own home.

data/

raw, processed, and feature-engineered datasets

src/

data, features, models, evaluation code

configs/

hyperparameters & pipeline settings as code

tests/

automated validation for code and data

notebooks/

exploration only — never production logic

dvc.yaml

reproducible pipeline definition

```
ml-project/
├── data/
│   ├── raw/
│   ├── processed/
│   └── features/
├── notebooks/
├── src/
│   ├── data/
│   ├── features/
│   ├── models/
│   └── evaluation/
├── configs/
├── tests/
├── mlruns/
├── dvc.yaml
└── requirements.txt
```

The MLOps Tool Landscape

Beyond MLflow and DVC — a broader ecosystem covers every stage of the ML lifecycle.



Source: InfluxData, "MLOps: A Comprehensive Guide to Machine Learning Operations"

Testing an ML System

Unit tests alone don't catch a model that's statistically wrong. ML systems need their own test pyramid.

FIG. 01

Data tests

Schema, ranges, nulls, and distribution checks on every incoming batch.

FIG. 03

Behavioral tests

Invariance and directional-expectation checks — e.g. changing an unrelated field shouldn't flip the prediction.

FIG. 02

Unit tests

Feature transforms and pipeline code, tested like any other software.

FIG. 04

Integration tests

The full pipeline end-to-end, including serving latency and contract checks.

Best Practices

Five habits that separate teams that ship ML reliably from teams that don't.

- 1

Treat Config as Code
Never hardcode hyperparameters — use versioned YAML config files.
- 2

Automate Data Quality Checks
Validate every batch with tools like Great Expectations before training.
- 3

Implement Model Monitoring
Detect data drift automatically and trigger retraining (e.g. Evidently).
- 4

Version Everything
Code in Git, data in DVC, models in MLflow Registry, infra in Terraform.
- 5

Shadow Mode Deployments
Run candidate models alongside production, compare before promoting.

Common Pitfalls & How to Avoid Them

Four mistakes that quietly sink ML projects — and the fix for each.

FIG. 01

Training/Serving Skew

PROBLEM: Features computed differently in training vs. production.

SOLUTION: Use a feature store or shared feature-engineering code for both.

FIG. 03

Ignoring Model Decay

PROBLEM: Accuracy quietly degrades as the world changes.

SOLUTION: Schedule regular retraining and monitor performance continuously.

FIG. 02

Data Leakage in CV

PROBLEM: Shuffling before split lets future data leak into training.

SOLUTION: Use time-based splits for temporal data.

FIG. 04

Over-Engineering Too Early

PROBLEM: Building Kubernetes clusters before validating model value.

SOLUTION: Follow crawl → walk → run: start simple, add complexity as value is proven.

06 · PITFALLS

Fixing a Common Pitfall: Time-Based Splits

Random shuffling is the single most common source of data leakage in time-series problems.

WRONG

```
# Wrong - future data leaks into training
X_train, X_test = train_test_split(
    X, test_size=0.2, shuffle=True)
```

RIGHT

```
# Right - split strictly by time
split_date = df['date'].quantile(0.8)
train = df[df['date'] < split_date]
test = df[df['date'] ≥ split_date]
```

An MLOps Maturity Model

Maturity is a gradient, not a switch. Most organizations move through these stages one capability at a time.



Implementing MLOps: A Roadmap

A high-level path organizations follow to stand up an MLOps practice.

- 1

Establish current state & objectives
Baseline today's deployment time and model accuracy; set concrete improvement goals.
- 2

Build the MLOps team
Blend data science, engineering, ops, and domain expertise.
- 3

Define data governance
Set standards for collection, storage, versioning, quality, and compliance.
- 4

Select tools & platforms
Weigh build vs. buy, team experience, and community support.
- 5

Automate the deployment pipeline
Start with CI/CD for testing, tracking, and promoting models.
- 6

Iterate and improve
MLOps is ongoing — revisit objectives and raise the bar over time.

CLOSING

Key Takeaways

FIG. 01

Version everything

Code, data, models, and config — all reconstructible.

FIG. 02

Automate the pipeline

From data validation through deployment, not just training.

FIG. 03

Monitor in production

Drift and decay are the default; watch for them continuously.

FIG. 04

Share ownership

MLOps succeeds as a team practice, not a handoff.

Thank you.

NEXT STEPS 1. Set up MLflow locally · 2. Version your first dataset with DVC · 3. Write one automated test · 4. Deploy to staging & monitor for a week